

Server Side Refactoring@AWS

by


Triton Ho@hypebeast.com



About hypebeast.com

- fashion website
- 3.7M monthly unique visitor
- Co-related website: hbx.com
- unpredictable traffic burst

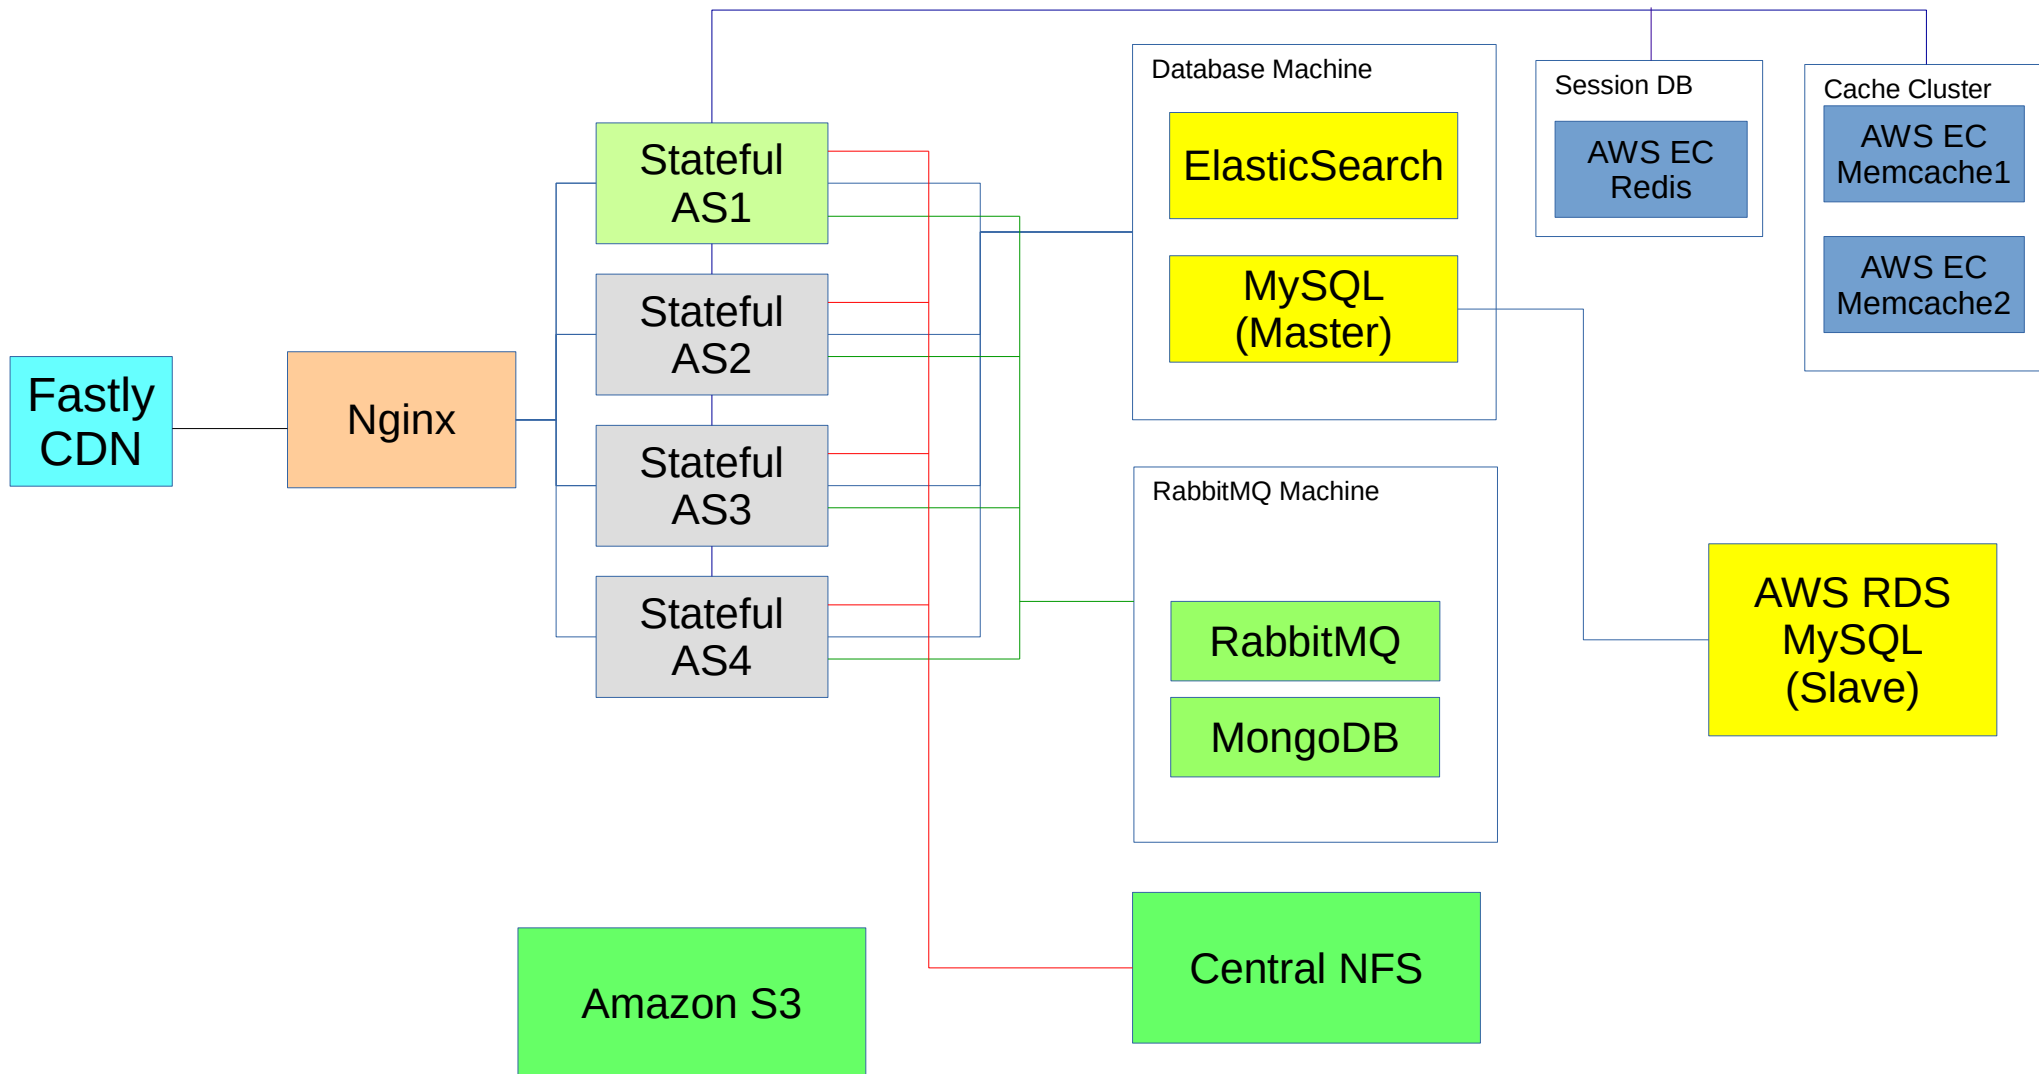
Technical challenges

- High traffic! High traffic! High traffic!
- Lack of manpower
 - 猫の手も借りたい 
- Brittle system
- Dogs in the office

Problems

- Multiple Single-Point-of-Failure
- Multiple performance bottleneck
- Hard for system upgrade
- Capacity Limit
- Incompatible components

Hypebeast.com@2015



Multiple Stage of Revamp

- Add kubernetes
- Migrate MySQL to Mariadb@RDS
- Migrate NFS to AWS EFS
- Migrate MongoDB, RabbitMQ to kubernetes
- Migrate ElasticSearch to AWS ES
- Migrate LoadBalancer to AWS ELB
- Bug fix on AS, become Stateless
- Migrate Memcache from AWS EC to AS localhost

kubernetes

- Motivation: too many system component
 - Example: Microsoft chatbot, multiple Wordpress
 - PHP 5.5, 5.6, 7.0.....
 - Node.JS 4.0, 6.0.....

Our Requirement on k8s

- As a platform to host various dockerized component
- Auto container restart after crash
- Store environment variable
- Easy service endpoint setup
- Allow Internal Service endpoint
- Deployment management

Why Not AWS ECS

- Troublesome and unmanagable
- Shattered documentation and components
 - AWS Route 53?
 - AWS CodeDeploy?
- Not a better solution than k8s

Deployment on k8s

- “kube deploy projectX” on slack
- HUBOT run deploy.sh
- `cd ~/projects/projectX && git pull`
- Build docker image
- Upload the docker image to AWS ECR
- Rolling deployment

Internal Service on k8s

- k8s has own Amazon VPC.
- Thus, VPC peering connection is needed.
- Add pcx-XXXXXXXXXX to your VPC Routing Table
- For each Internal Service
 - Create k8s services(type = NodePort).
 - Create a AWS Internal ELB
 - Add the Internal ELB to k8s Auto Scaling Group

Afterthought on k8s + docker

- Costed one month to study + minesweeping + deploy.
- 10.0.0.0/8 is hardcoded as pod's IP.
 - Collides with AWS RDS default IP range.
- Default instance type m3.medium is too small
- For new service:
 - ≥ 1 days to prepare the DockerFile
 - 2-3 hours to setup on k8s
- Private library repo caused great trouble in DockerFile

Mariadb@RDS

- Motivation:
 - Mariadb is more advanced
 - Replication/Failover is not easy and error-prone
 - Setup is easy, but automation is hard.
- Self Hosted instance is hard for applying patch, instance scaling

Near-Zero Migration To RDS

- mysqldump the production mysql
- Create the RDS instance(with Multi-AZ) and then import the data
- Setup the replication from mysql to mariadb
 - mysql.rds_start_replication on RDS
- Switch the production connection from mySQL to mariadb
- Stop the replication
 - mysql.rds_stop_replication on RDS

Afterthought on RDS

- Highly vendor lock-in.
 - No external replication
- Multi-AZ
 - Robbery, 2x price but no readonly-slave
 - After instance-type change, the availability-zone change too
 - EXTRA 1MS DOES MATTER FOR BAD CODE
- No c4 instance, bad for MySQL and mariadb
- In config, no CPUNum variable

Amazon EFS

- Motivation:
 - Need a flexible size NFS
 - Need a global storage for wordpress

Afterthought on AWS EFS

- Just a NFS
- On Ubuntu 16.04.1, failed to automount
- No monitoring
- Poor performance for hosting php code
 - Any one particular file can have up to 87 locks across all users of the file system.


Mongodb+RabbitMQ@k8s

- Motivation:
 - Used by side project only, low usage
 - Switch off the EC2 instance

Afterthought on k8s

- Amazon internal-ELB has idle timeout 60s
 - Our mongodb / MQ connection raise exception every 60 sec.
 - Increase the idle time to 3600 sec
- EBS is used for data storage.
- Why not AWS SQS:
 - We need delayed Messaging.

AWS ES

- Motivation:
 - Lack of Manpower
 - 猫の手も借りたい 
 - Scalability and HA

Afterthought on AWS ES

- WTF, no security group and security group
- For hypebeast.com, t2.medium is better

From nginx to AWS ELB

- Motivation:
 - Central Nginx become bottleneck
 - Central Nginx is the single point of failure
- Original:
 - Two nginx tier
 - One central nginx for LB(ip_hash) and url rewrite
 - Nginx on each application server for PHP

Afterthought on AWS ELB

- Sticky loadBalancer will override X-Forwarded-Proto
 - Serious problem for CDN
- Non Sticky loadBalancer raise bug in our code.
- Shit.....

Revamp on AS

- Motivation:
 - Make the AS become stateless
 - Increase performance by upgrade OS version and PHP version
 - Reduce Network latency by using localhost memcache

Afterthought on AWS ElasticCache

- Just waste your money
- Application server is CPU-bound, Cache server is RAM-bound
 - Combining them is perfect
- Using localhost memcache, the avg time on cache decreased from 30ms to 8ms

Conclusion

- AWS is good for startup to save manpower
- But AWS feature is not enough for advanced usage
 - Especially you have technical debt
- If you just need VM, gCloud may be much better.