# Effective Data Analysis and ISO 2008 SQL Features by Triton Ho

## **Today Contents**

- Introductory Examples
- Basic Concepts on Reporting
- ISO SQL:2008 standard

Find out the eldest child for each family
 Select \* from students s1
 where not exists (
 select 1 from students s2
 where s1.parent\_id = s2.parent\_id
 and s2.age > s1.age

Find out the products stored in multiple warehouses

```
Select product_id
from inventories t1
group by product_id
having count(distinct warehouse_id) > 1
```

Just-in-time calculation of player's action point select id, least(100, candy + trunc(extract(epoch from current\_timestamp - last\_candy\_time) / 60 / 8)
 ) as candy from players

#### Basic Concepts on Reporting

# **Basic Concepts on Reporting**

- Avoid select \*
- Cluster your query
- SQL Join is not slow
- Seq. Scan is not bad
- Caching and Seq. Scan
- Introduction to Query Optimizer
- CTAS with Temp. Table

#### Avoid select \*

- Always explicitly list out your columns
- Avoid wasteful network IO
- Maybe better execution plan
  - Index-only-scan?

#### Stupid Example 1

```
idArray := []int{1, 43, 76, 757}
for _, id := range idArray {
    q := `select * from products where id = ?`
    db.RunQuery(q, id)
    .....
}
```

# Cluster your Query

- Network overhead costs time
- SQL parsing and optimization takes time
- Computation should be close to the data
  - If necessary, move your data
- In Summary: your number of query should be independent of your data size.

#### Good Example 1

```
idArray := [int{1, 43, 76, 757}]
temp := []string{}
for , := range idArray {
  temp = append(temp, `? `)
q := `select * from products ` +
`where id in (` + strings.Join(temp, `,`) + `)`
db.RunQuery(q, id...)
```

# Brainstorm: Dijkstra's algorithm

- If the dataset is small, query all edges data into memory.
- Write stored procedure and run the algorithm in RDBMS directly
- Cache the edges data in application tier

#### SQL Join is not slow

 Example: we want to display the flight details in additional to the flight ticket.

```
select * from flight_ticket ft
inner join flight f on ft.flight_id = f.id
where ft.id = @ticket_id
```

## Explain the SQL

- ft.id = @ticket\_id is filtering on PK
  - flight\_ticket has one record only
- f.id is the PK of flight, and thus has index
  - It will be the inner loop of index nested loop join

# Index nested loop join

- Step1: flitering on flight\_ticket by ft.id =
   @ticket\_id
- Step2:

```
- foreach candidate ft in flight_ticket {
    f := flight.getRecordByIndex(ft.flight_id)
    if f exists, add <f, ft> into result
}
```

## Seq. Scan is not bad

- Every man knows:
  - Copying one 10GB movie file is faster than copying one thousand 1MB jpeg files
- Usually, RDBMS ignore secondary index(i.e. non-PK)

#### Seq. Scan with PK

- Applicable to clustered index only
  - Oracle IOT, Mariadb(innodb) table
- Sometimes, you want to get records of a period.
  - e.g. create\_time between '2016-01-01' and '2017-12-31'
- Extra predicate: "id between XXX and YYY" may help to reduce the seq. scan range

# Caching and Seq. Scan

- Seq. Scan on 5GB table ALWAYS get cache miss on 4GB RAM LRU cache.
  - Oracle / Postgresql on purposely skip caching on Seq. Scan
- For reporting, RAM is useful for
  - storing intermediate dataset
  - Caching random access(e.g. index join)

# Introduction to Query Optimizer

- In oracle / postgreSQL, the following SQL are SAME
  - Select child.id from child inner join parent on child.parent\_id = parent.id where parent.name = 'TritonHo'
  - Select child.id from child where child.parent\_id = (select id from parent where parent.name = 'TritonHo')

```
    Select child.id from child where exists
        (
            select 1 from parent
            where child.parent_id = parent.id and parent.name = 'TritonHo'
        )
```

# SQL Optimizer

- Good SQL optimizer ignores
  - joining sequence
  - Location of predicate (i.e. in where clause or as a joining condition)
- Good SQL optimizer transforms sub-query into equivalent join statement
- Optimizer Ranking
  - Oracle > Postgresql >>> Mariadb >= MySQL

# SQL Optimizer(cont'd)

- The difference between genius and stupidity is that genius has its limits.
- Runtime of SQL optimizer is O(2<sup>n</sup>)
   n = # of table involved
- If n is too large, optimizer will give up and use heuristic
  - My experience: >50 lines is dangerous for reporting

#### Temp. Table

- on transaction rollback / commit, automatic deleted
- Data in Temp. Table will not be shared between transactions
- Much faster than normal tables.
- You can perform insert / update / delete

#### Temp. Table and CTAS

- CTAS = create table as select ......
- Using CTAS, you can store intermediate dataset easily
- Thus, you can break one complex query into multiple simple query
- Example: find the 2<sup>nd</sup> child in all families

## Temp. Table example

Step 1: find the 1<sup>st</sup> child and save the result

```
create temp table temp1 as
Select id from students s1
where not exists (
    select 1 from students s2
    where s1.parent_id = s2.parent_id
    and s2.age > s1.age
)
```

## Temp. Table example

Step 2: minus the 1st child and save the result

```
create temp table t2 as
Select id from students s1
where s1.id not in (
select temp1.id from temp1
)
```

## Temp. Table example

Step 3: get the oldest child in the dataset

```
Select id from temp2 a
where not exists (
select 1 from temp2 b
where a.parent_id = b.parent_id
and b.age > a.age
)
```

ISO SQL:2008 standard

#### Great Milestone of ISO SQL

- SQL:1999
  - With-Clause
- SQL:2003
  - Window function
  - CTAS
- SQL:2008
  - Case-when expression
- MySQL / MariaDB
  - SQL:1992

#### Without with-clause

```
Select * from (
    Select * from (
    SELECT class id, student id, parent id
    rank() over (partition by class_id order by score desc) as ranking
    FROM students
    where students.group_id = 2
    ) t1
    where t1.ranking <= 3
) wtf1
where exists (
    Select 1 from (
    SELECT class_id, student_id, parent_id
    rank() over (partition by class id order by score desc) as ranking
    FROM students
    where students.group id = 2
    ) t2
    where t2.ranking <= 3
    and wtf1.student id != t2.student id and wtf1.parent id = t2.parent id
```

#### With-clause

- Not support by MySQL / MariaDB
- Allow just-in-time view declaration
- Implicit materialization
  - Affect execution plan, both good and bad
  - Need manual predicate pushing

# With-clause example

```
With top3students as (
   select * from (
        SELECT class id, student id, parent id
        rank() over (partition by class id order by score desc) as ranking
        FROM students
        where group id = 2
   ) t1
   where t1.ranking <= 3
Select * from top3students t1
where exists (
    Select 1 from top3students t2
   where t1.student id != t2.student id and t1.parent id = t2.parent id
```

# predicate pushing

- "push" your predicate to inner layer(i.e. subquery)
- Automatically done by query optimizer
  - Again: Oracle > Postgresql >>> Mariadb >= MySQL
- Blocked by
  - Materialization(e.g. with-clause)
  - Aggregation(COUNT, SUM, AVG)
  - Window function

## Bad example

```
With top3students as (
    select * from (
        SELECT class id, student id, parent id
        rank() over (partition by class_id order by score desc) as ranking
        FROM students
    ) t1
    where t1.ranking <= 3
Select * from top3students t1
where exists (
    Select 1 from top3students t2
    where t1.student_id != t2.student_id and t1.parent_id = t2.parent_id
    and t1.group id = t2.group id
and t1.group id = 2
```

#### Crazy example

```
Select group_id, avg(score) from students group by group_id having group_id = 2
```

#### Window-function

Before SQL:2003, pagination example:

```
select * from students where ..... order by id limit 50
```

- Pagination is not standardized
  - Oracle: rownum
  - Postgresql / Mariadb: limit and offset
  - MSSQL: Top
- Window function is standardized and supported by all SQL:2003 compliant DB
  - Once again: Sorry for Mariadb users

Pagination by window function

```
select * from (
    select *,
    row_number() over (order by score desc) as row_num
    from students
    where .....
) t
where t.row_num between 11 and 20
```

Top x item in each group

```
Select * from (
    SELECT class_id, student_id,
    rank() over (partition by class_id order by score desc) as
ranking
    FROM students
) t
where t.ranking <= 3
```

- Much higher performance
  - Single pass on the table
  - No joining, no union

Facebook-like thread

```
WITH RECURSIVE message with replies as (
 select * from (
  SELECT to char(row number() over(order by create time asc), 'fm0000000') as seq, "::text as parent seq,
null::bigint as row num, *
  FROM filtered messages
  where .....
  order by create time asc
  offset? Limit?
 ) t1
 UNION ALL
 select parent seq || '-' || to char(row num, 'fm0000000') as seq, * from (
  SELECT p.seg as parent seg, row number() over(partition by r.parent id order by r.create time asc) as row num,
  FROM message with replies p
  inner join filtered messages r on p.id = r.parent id
 ) t2
 where row num <=?
), filtered messages as (
 select * from messages where ......
SELECT * from message with replies
order by seq
```

As a syntax sugar for aggregation

```
SELECT class_id, student_id, score avg(score) over (partition by class_id) as avg_score FROM students
```

- Traditional group-by:
  - Select class\_id, student\_id, score, t1.avg\_score from students s1 inner join (
     select class\_id, avg(score) as avg\_score from students s2
     group by class\_id
     ) t2 on s1.class\_id = t2.class\_id

# Case-when expression

- Allow you perform "switch"
- Similar function: coalesce

## Case-when Example

 You want to get girls in ascending order, and then boys in descending order

```
Select * from students
order by gender,
(
    case
    when gender = 'female' then height
    when gender = 'male' then height * -1
    end
)
```

#### Conclusion

- SQL is a declarative language
  - Elegant and Powerful
  - Easy to write, hard to master
- · Unless debugging, optimizer hints is bad idea

# Some bad example

```
Select ......
from tableA
left join tableB on tableA.XXX = tableB.YYY
where tableB.ZZZ = ?
```

#### End