

SQL Reporting 淺談 v2.0

by

Triton Ho



今天內容

- 簡單例子
- 基礎觀念
- ISO SQL:2008 標準

前言

- 沒人叫你把一切數據分析都用 SQL 來跑
 - TB 級數據面前，SQL 有高度局限性
 - Python 和 R 的統計功能非常強大，而且能用 GPU
- 但是……
 - 一般中小型公司的數據量 < 100GB
 - SQL 是宣告式語言，語法簡單短小
 - 用 Python 和 R 來做分析時，還是要用 SQL 先拿資料出來

簡單例子 1

- 找出每個家庭的最年長的孩子

```
Select * from residents r1
where not exists (
    select 1 from residents r2
    where r1.parent_id = r2.parent_id
    and r2.age > r1.age
)
```

簡單例子 2

- 找出在多於一個倉庫中有存貨的產品

```
Select product_id  
from inventories t1  
where amount > 0  
group by product_id  
having count(distinct warehouse_id) > 1
```

簡單例子 3

- 即時計算玩家的行動點數

```
select id, least(100, candy +  
trunc(extract(epoch from current_timestamp -  
last_candy_time) / 60 / 8)  
) as candy  
from players
```

簡單例子 4

- 把用戶的身高，以每 5cm 為單位，統計用戶人數

```
select coalesce(t1.counter, 0) as counter
from generate_series(1,20) t2
left join (
    select WIDTH_BUCKET(height, 140, 200, 12) as
basket_id,
    count(*) as counter
    from users
    group by basket_id
) t1 on t1.basket_id = t2
order by t2
```

簡單例子 5

- 以每 5% 作解像度，計算用戶的身高百分位數 (percentile)

```
select percentile_cont(array(SELECT j/20.0  
FROM generate_series(1,20) j))  
within group (order by height) as  
height_percentile  
from users
```


基礎觀念

基礎觀念

- 避免 Select *
- 避免散碎的 Query
- SQL Join 並不慢
- Seq. Scan 不一定是壞事
- Caching 和 Seq. Scan
- Query Optimizer 淺談
- 善用 CTAS 和 Temp table

避免 Select *

- 請盡量把你需要的 column 明確地寫出來
- 別浪費 Network IO
 - 特別是你的資料量是 GB 級時
- 有可能會有更好的 execution plan
 - 例子：Index-only scan

笨例子 1

```
idArray := []int{1, 43, 76, 757}
for _, id := range idArray {
    q := `select * from products where id = ?`
    db.RunQuery(q, id)
    .....
}
```

避免散碎的 Query

- 網絡通訊會吃掉時間
- SQL Parsing 和優化也吃時間
- 重要觀念：
 - 數據運算應該接近數據所在的位置
(Computation should be close to the data)
 - 另一個說法：如果必需，先把資料移動，然後才做運算
- 簡單來說：
 - 你的 Query 次數應該跟數據量無關

好例子 1

```
idArray := []int{1, 43, 76, 757}
```

```
temp := []string{}
```

```
for _, _ := range idArray {  
    temp = append(temp, `?`)  
}
```

```
q := `select * from products ` +  
`where id in (` + strings.Join(temp, `,`) + `)`  
db.RunQuery(q, id...)
```

思考：Dijkstra's 演算法

- 如果數據量不大，乾脆把整個 graph 丟到 application 記憶體吧 ~
- 數據量很大時：
 - 寫 stored procedure(pl/pgsql) ，在資料庫上面跑算法 ~
 - 每次拿路徑 (edge) 資料時，除了目標節點 (node) ，還把相鄰的 (node) 也拿出來，然後做好 Caching

SQL Join 並不慢 1

- 例子：在顯示機票資料時，我們要顯示相關的航次資料

```
select * from flight_ticket ft  
inner join flight f on ft.flight_id = f.id  
where ft.id = @ticket_id
```


SQL Join 並不慢 2

- `ft.id = @ticket_id` 是以 Primary key 來查詢
 - 即是說：這個最多只有一個結果
- `f.id` 是 flight 這個表的 PK，所以會有對應的 index
 - 所以，這個 `f.id` 會放在 **index nested loop join** 的內圈

Index nested loop join

- Step1: `ft.id = @ticket_id` , 對 `flight_ticket` 做搜查
- Step2:
 - foreach candidate `ft` in `flight_ticket` {
 `f := flight.getRecordByIndex(ft.flight_id)`
 if `f` exists, add `<f, ft>` into result
}

Seq. Scan 不一定是壞事

- 每個人都知道：
 - 抄一個 10GB 檔案，比一千個 1MB 檔案更快
- 在報表類應用， RDBMS 一般會無視 secondary index(i.e. non-PK)

Seq. Scan with PK

- 只適用於 clustered index
 - Oracle Index-organized-table
 - Mariadb(innodb) table
- 如果你想拿某時段的資料
 - create_time between '2016-01-01' and '2017-12-31'
 - 做年度報表經常用上
- Extra predicate: “id between XXX and YYY” 可能有幫忙

•Caching 和 Seq. Scan

- 對 5GB table 作 Seq. Scan 時，如果你資料庫只有 4GB RAM
 - Mariadb 是用 LRU cache 的
這情況下 cache hit rate = 0
 - Oracle / Postgresql 不會對 Seq. Scan 作 caching
- 在報表類 Query 時，RAM 是用作：
 - 暫時儲存中間的運算結果
 - 對 Random IO 作 Caching
(例子：index join)

Query Optimizer 淺談

- 在 oracle / postgresSQL ，以下 Query 效能相同
 - Select child.id from child
inner join parent on child.parent_id = parent.id
where parent.name = 'TritonHo'
 - Select child.id from child
where child.parent_id =
(select id from parent where parent.name = 'TritonHo')
 - Select child.id from child
where exists
(
 select 1 from parent
 where child.parent_id = parent.id and parent.name = 'TritonHo'
)

SQL Optimizer

- 再次強調：SQL 是宣告式語言
- 好的 SQL optimizer 讓使用者不用思考
 - joining sequence
 - predicate location
- 好的 SQL optimizer 會把 sub-query 改成 join
- Optimizer Ranking
 - Oracle > Postgresql >>> Mariadb >= MySQL

SQL Optimizer (續)

- 天材有其極限，笨蛋沒有
- SQL optimizer 的工作時間 $\sim O(2^n)$
 $n = \# \text{ of table}$
- 如果 n 太大，optimizer 會改用猜猜樂 ~
猜猜樂的好聽說法：heuristic / 先進預測
 - 個人看法：報表類 SQL，不建議超過 50 行

Temp. Table

- 在 TX 結束時，其內資料全部被刪除
- Temp. Table 內資料無法被其他 connection 看到
- 效能比一般的 Table 快很多
 - 有些資料庫支援以 RAM 作 Temp. Table 的儲存空間
- 能做 insert / update / delete

Temp. Table 和 CTAS

- CTAS = create table as select
- 使用 CTAS ，你能把中間的結果用 Temp. Table 暫存起來
- 所以，你可以把一個複雜的 SQL 變成數個簡單的
- 以下例子：要找到每個家庭的第二個孩子

Temp. Table 例子 1

- Step 1: 找到最年長的孩子，然後把結果存到 temp1

```
create temp table temp1 as
Select id from residents r1
where not exists (
    select 1 from residents r2
    where r1.parent_id = r2.parent_id
    and r2.age > r1.age
)
```

Temp. Table 例子 2

- Step 2: 把本來的 students ，扣掉 temp1 的記錄，再存到 temp2

```
create temp table temp2 as  
Select id from residents r1  
where r1.id not in (  
    select temp1.id from temp1  
)
```

Temp. Table 例子 3

- Step 3: 從 temp2 中找到其內每個家庭最年長的孩子

```
Select id from temp2 a
  where not exists (
    select 1 from temp2 b
    where a.parent_id = b.parent_id
    and b.age > a.age
  )
```

ISO SQL:2008 標準

ISO SQL 里程碑

- SQL:1999
 - With-Clause
- SQL:2003
 - Window function
 - CTAS
- SQL:2008
 - Case-when expression
- MySQL 5.7 / MariaDB 10.1 目前支援度：
 - SQL:1992

沒有 with-clause 時

```
Select * from (  
  Select * from (  
    SELECT class_id, student_id, parent_id  
    rank() over (partition by class_id order by score desc) as ranking  
    FROM students  
  where students.group_id = 2  
  ) t1  
  where t1.ranking <= 3  
) wtf1  
where exists (  
  Select 1 from (  
    SELECT class_id, student_id, parent_id  
    rank() over (partition by class_id order by score desc) as ranking  
    FROM students  
  where students.group_id = 2  
  ) t2  
  where t2.ranking <= 3  
  and wtf1.student_id != t2.student_id and wtf1.parent_id = t2.parent_id  
)
```


With-clause

- 再說一次：
 - MariaDB 使用者，你可以今天後回家哭哭了 ~
- 讓你建立 Just-in-time ，單次性使用的 view
- 自動（強制性的） materialization
 - 會影響到 Execution plan ，是好事也是壞事
 - 需要手動 predicate pushing

With-clause 例子

```
With top3students as (  
    select * from (  
        SELECT class_id, student_id, parent_id  
        rank() over (partition by class_id order by score desc) as ranking  
        FROM students  
        where group_id = 2  
    ) t1  
    where t1.ranking <= 3  
)  
Select * from top3students t1  
where exists (  
    Select 1 from top3students t2  
    where t1.student_id != t2.student_id and t1.parent_id = t2.parent_id  
)
```

predicate pushing

- 把 predicate 從外層推送到內層 (i.e. subquery)
- 正常來說， query optimizer 會替你自動完成
 - 還是一句： Oracle > Postgresql >>> Mariadb >= MySQL
- query optimizer 無法 predicate pushing 例子：
 - Materialization(e.g. with-clause)
 - Aggregation(COUNT, SUM, AVG)
 - Window function

壞例子

```
With top3students as (  
    select * from (  
        SELECT class_id, student_id, parent_id  
        rank() over (partition by class_id order by score desc) as ranking  
        FROM students  
    ) t1  
    where t1.ranking <= 3  
)  
Select * from top3students t1  
where exists (  
    Select 1 from top3students t2  
    where t1.student_id != t2.student_id and t1.parent_id = t2.parent_id  
    and t1.group_id = t2.group_id  
)  
and t1.group_id = 2
```

瘋狂例子

```
Select group_id, avg(score)
from students
group by group_id
having group_id = 2
```

Window-function

- 在 SQL:2003 之前，如果要做分頁 (pagination) :

```
select * from students  
where ..... order by id limit 50
```
- 沒有統一標準
 - Oracle: rownum
 - Postgresql / Mariadb: limit and offset
 - MSSQL: Top
- Window function 是統一標準，所有 ISO SQL:2003 compliant 都支援

Window-function 例子 1

- 分頁功能

```
select * from (  
    select *,  
    row_number() over (order by score desc) as row_num  
    from students  
    where .....  
) t  
where t.row_num between 11 and 20
```

Window-function 例子 2

- 想拿到每班別的首三名學生：

```
Select * from (  
    SELECT class_id, student_id,  
    rank() over (partition by class_id order by score desc) as ranking  
    FROM students  
    ) t  
where t.ranking <= 3
```

- 遠遠優勝的效能
 - table 只需要被掃描一次
 - 不需要 joining ，也不需要 union

Window-function 例子 3

- Facebook-like thread

```
WITH RECURSIVE message_with_replies as (  
  select * from (  
    SELECT to_char(row_number() over(order by create_time asc), 'fm0000000') as seq, '::text' as parent_seq,  
    null::bigint as row_num, *  
    FROM filtered_messages  
    where .....  
    order by create_time asc  
    offset ? Limit ?  
  ) t1  
  UNION ALL  
  select parent_seq || '-' || to_char(row_num, 'fm0000000') as seq, * from (  
    SELECT p.seq as parent_seq, row_number() over(partition by r.parent_id order by r.create_time asc) as row_num,  
    r.*  
    FROM message_with_replies p  
    inner join filtered_messages r on p.id = r.parent_id  
  ) t2  
  where row_num <= ?  
, filtered_messages as (  
  select * from messages where .....  
)  
SELECT * from message_with_replies  
order by seq
```

Window-function 例子 4

- 作為 syntax sugar

```
SELECT class_id, student_id, score  
avg(score) over (partition by class_id) as avg_score  
FROM students
```

- 傳統的 group-by:

- Select class_id, student_id, score, t1.avg_score
from students s1 inner join (
 select class_id, avg(score) as avg_score
 from students s2
 group by class_id
) t2 on s1.class_id = t2.class_id

Case-when

- 容許你在 SQL 內做到 switch
- 相似例子：coalesce

Case-when 例子

- 在拍班級照片時，你想左邊全是女生，右邊全是男生
學生高度是兩邊向中央升高
- ```
Select * from students
order by gender,
(
 case
 when gender = 'female' then height
 when gender = 'male' then height * -1
 end
)
```

# 結語

- SQL 是宣告式語言
  - 優美而且功能強大
  - 在 GB 級數據量，能快速寫好
  - 易學難精 > < "
- MySQL / Mariadb 最大問題：
  - 只去做搶眼球的門面功夫
  - 但是卻不肯練好內功

End